

```
In [55]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import warnings
warnings.filterwarnings('ignore')
```

```
In [56]: dataframe = pd.read_csv("winequality.csv")
dataframe.head()
```

Out[56]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qualit
0	white	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0.45	8.8	
1	white	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0.49	9.5	
2	white	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0.44	10.1	
3	white	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	
4	white	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	

```
In [57]: dataframe.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   type                                6497 non-null   object
1   fixed acidity                       6487 non-null   float64
2   volatile acidity                    6489 non-null   float64
3   citric acid                         6494 non-null   float64
4   residual sugar                      6495 non-null   float64
5   chlorides                           6495 non-null   float64
6   free sulfur dioxide                 6497 non-null   float64
7   total sulfur dioxide                6497 non-null   float64
8   density                             6497 non-null   float64
9   pH                                  6488 non-null   float64
10  sulphates                           6493 non-null   float64
11  alcohol                             6497 non-null   float64
12  quality                             6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

```
In [58]: dataframe.describe().T
```

Out[58]:

	count	mean	std	min	25%	50%	75%	max
fixed acidity	6487.0	7.216579	1.296750	3.80000	6.40000	7.00000	7.70000	15.90000
volatile acidity	6489.0	0.339691	0.164649	0.08000	0.23000	0.29000	0.40000	1.58000
citric acid	6494.0	0.318722	0.145265	0.00000	0.25000	0.31000	0.39000	1.66000
residual sugar	6495.0	5.444326	4.758125	0.60000	1.80000	3.00000	8.10000	65.80000

<b>chlorides</b>	6495.0	0.056042	0.035036	0.00900	0.03800	0.04700	0.06500	0.61100
<b>free sulfur dioxide</b>	6497.0	30.525319	17.749400	1.00000	17.00000	29.00000	41.00000	289.00000
<b>total sulfur dioxide</b>	6497.0	115.744574	56.521855	6.00000	77.00000	118.00000	156.00000	440.00000
<b>density</b>	6497.0	0.994697	0.002999	0.98711	0.99234	0.99489	0.99699	1.03898
<b>pH</b>	6488.0	3.218395	0.160748	2.72000	3.11000	3.21000	3.32000	4.01000
<b>sulphates</b>	6493.0	0.531215	0.148814	0.22000	0.43000	0.51000	0.60000	2.00000
<b>alcohol</b>	6497.0	10.491801	1.192712	8.00000	9.50000	10.30000	11.30000	14.90000
<b>quality</b>	6497.0	5.818378	0.873255	3.00000	5.00000	6.00000	6.00000	9.00000

```
In [59]: dataframe.isnull().sum()
```

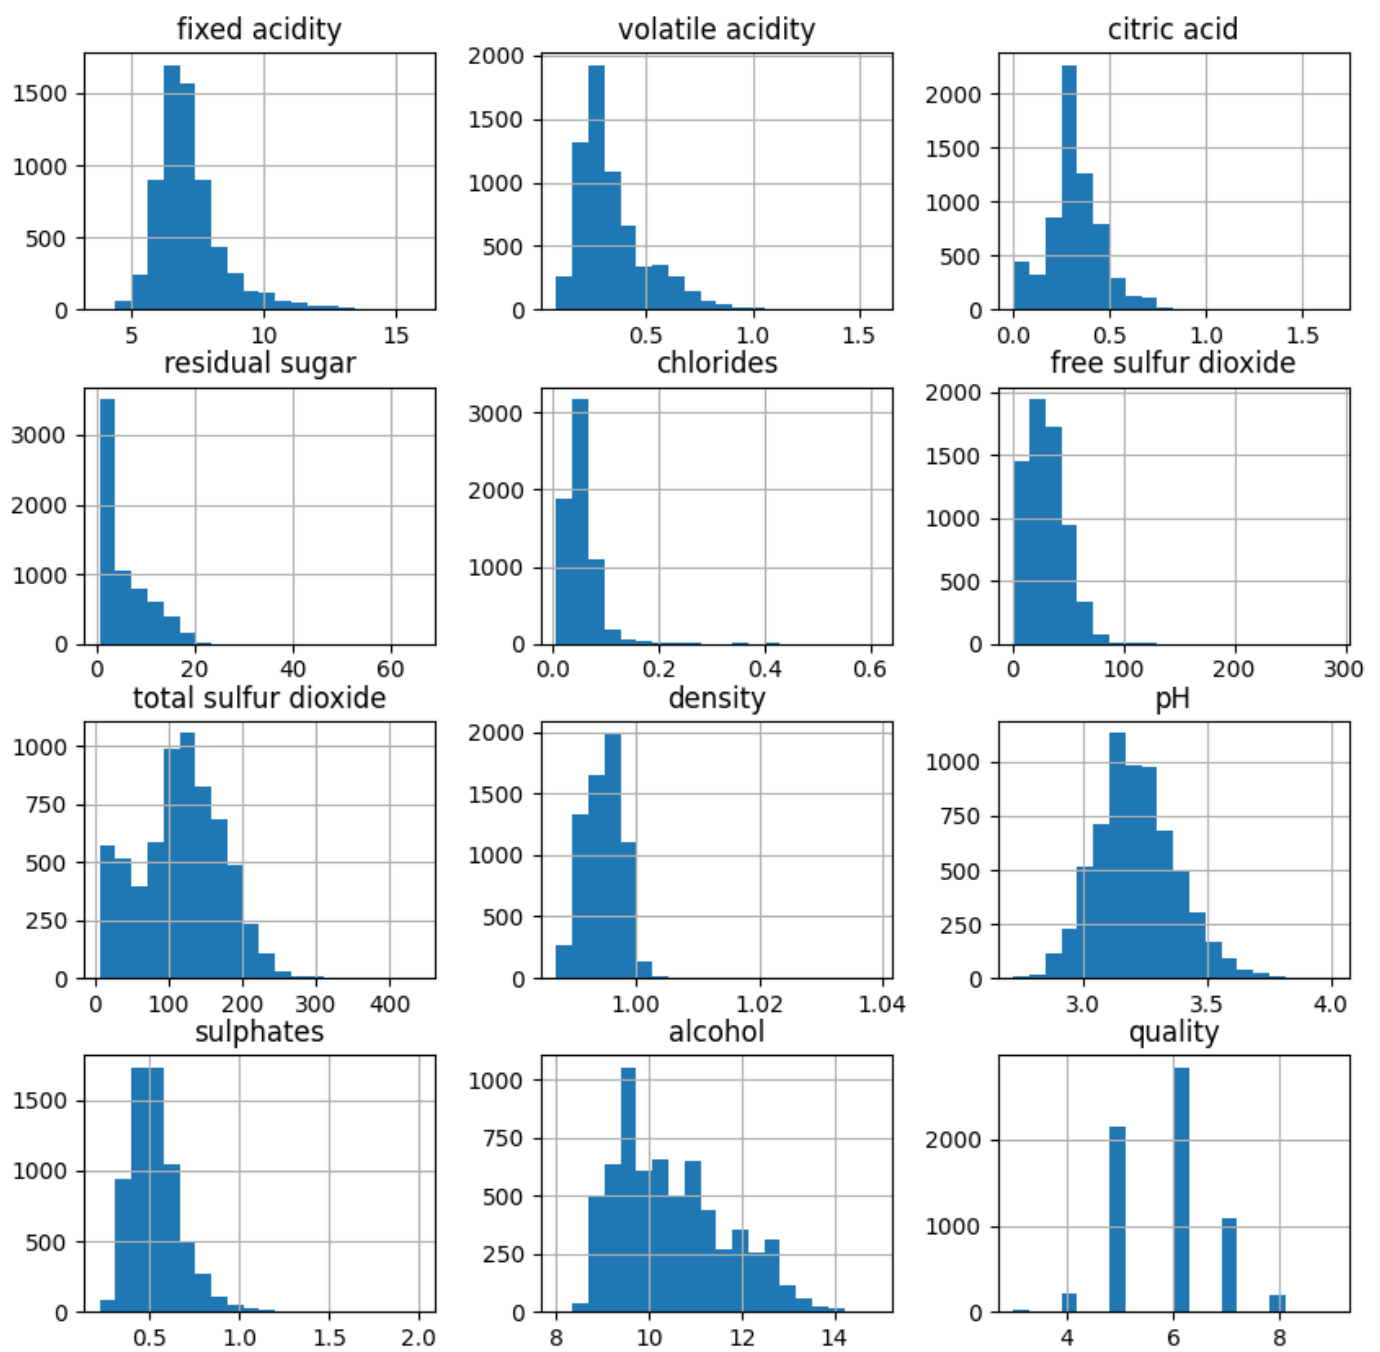
```
Out[59]: type                0
fixed acidity            10
volatile acidity         8
citric acid              3
residual sugar           2
chlorides                2
free sulfur dioxide      0
total sulfur dioxide     0
density                 0
pH                      9
sulphates               4
alcohol                 0
quality                 0
dtype: int64
```

```
In [60]: for col in dataframe.columns:
          if dataframe[col].isnull().sum() > 0:
              dataframe[col] = dataframe[col].fillna(dataframe[col].mean())


dataframe.isnull().sum().sum()
```

```
Out[60]: 0
```


```
In [61]: dataframe.hist(bins = 20, figsize = (10, 10))
plt.show()
```



```
In [62]: plt.bar(dataframe['quality'], dataframe['alcohol'])
plt.xlabel('quality')
plt.ylabel('alcohol')
plt.show()
```



```
In [63]: plt.figure(figsize = (12, 12))
range_except_date = dataframe.loc[:, dataframe.columns != 'type']
sb.heatmap(range_except_date.corr() > 0.7, annot = True, cbar = False)
plt.show()
```



```
In [64]: dataframe = dataframe.drop('total sulfur dioxide', axis = 1)
```

```
In [65]: dataframe = dataframe.drop('free sulfur dioxide', axis = 1)
```

```
In [66]: dataframe.head()
```

Out[66]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	density	pH	sulphates	alcohol	quality
0	white	7.0	0.27	0.36	20.7	0.045	1.0010	3.00	0.45	8.8	6
1	white	6.3	0.30	0.34	1.6	0.049	0.9940	3.30	0.49	9.5	6
2	white	8.1	0.28	0.40	6.9	0.050	0.9951	3.26	0.44	10.1	6
3	white	7.2	0.23	0.32	8.5	0.058	0.9956	3.19	0.40	9.9	6
4	white	7.2	0.23	0.32	8.5	0.058	0.9956	3.19	0.40	9.9	6

```
In [67]: dataframe['best quality'] = [1 if x > 5 else 0 for x in dataframe.quality]
```

```
In [68]: dataframe.head()
```

```
Out[68]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	density	pH	sulphates	alcohol	quality	best quality
0	white	7.0	0.27	0.36	20.7	0.045	1.0010	3.00	0.45	8.8	6	1
1	white	6.3	0.30	0.34	1.6	0.049	0.9940	3.30	0.49	9.5	6	1
2	white	8.1	0.28	0.40	6.9	0.050	0.9951	3.26	0.44	10.1	6	1
3	white	7.2	0.23	0.32	8.5	0.058	0.9956	3.19	0.40	9.9	6	1
4	white	7.2	0.23	0.32	8.5	0.058	0.9956	3.19	0.40	9.9	6	1

```
In [69]: dataframe.replace({'white': 1, 'red': 0}, inplace = True)
```

```
In [70]: dataframe.head()
```

```
Out[70]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	density	pH	sulphates	alcohol	quality	best quality
0	1	7.0	0.27	0.36	20.7	0.045	1.0010	3.00	0.45	8.8	6	1
1	1	6.3	0.30	0.34	1.6	0.049	0.9940	3.30	0.49	9.5	6	1
2	1	8.1	0.28	0.40	6.9	0.050	0.9951	3.26	0.44	10.1	6	1
3	1	7.2	0.23	0.32	8.5	0.058	0.9956	3.19	0.40	9.9	6	1
4	1	7.2	0.23	0.32	8.5	0.058	0.9956	3.19	0.40	9.9	6	1

```
In [71]: features = dataframe.drop(['quality', 'best quality'], axis = 1)
target = dataframe['best quality']

xtrain, xtest, ytrain, ytest = train_test_split(features, target, test_size = 0.2, random_state = 42)

xtrain.shape, xtest.shape
```

```
Out[71]: ((5197, 10), (1300, 10))
```

```
In [72]: norm = MinMaxScaler()
xtrain = norm.fit_transform(xtrain)
xtest = norm.transform(xtest)
```

```
In [73]: models = [LogisticRegression(), XGBClassifier(), SVC(kernel = 'rbf')]
for i in range(3):
    models[i].fit(xtrain, ytrain)
    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(ytrain, models[i].predict(xtrain)))
    print('Validation Accuracy : ', metrics.roc_auc_score(ytest, models[i].predict(xtest)))
    print()
```

```
LogisticRegression() :
Training Accuracy : 0.7019886368161423
Validation Accuracy : 0.7019518599115251
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
```

```

        multi_strategy=None, n_estimators=None, n_jobs=None,
        num_parallel_tree=None, random_state=None, ...) :
Training Accuracy : 0.969661860064318
Validation Accuracy : 0.788761353071587


SVC() :
Training Accuracy : 0.706251343942582
Validation Accuracy : 0.7030245910850932

```

```

In [74]: clf = SVC(random_state = 0)
         clf.fit(xtrain, ytrain)
         predictions = clf.predict(xtest)
         cm = confusion_matrix(ytest, predictions, labels = clf.classes_)
         disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels = clf.classes_)
         disp.plot()
         plt.show()

```



```

In [75]: print(metrics.classification_report(ytest, models[1].predict(xtest)))

```

	precision	recall	f1-score	support
0	0.76	0.70	0.73	474
1	0.84	0.87	0.85	826
accuracy			0.81	1300
macro avg	0.80	0.79	0.79	1300
weighted avg	0.81	0.81	0.81	1300